

Mastering Large-Size HTTP Requests in the Modern Web

Cherie Hsieh

About Me



跨領域工作

網路行銷 & PM

Web 前端工程師

Web Golang 後端工程師

韌體工程師

TSMC SRE



Golang Taiwan Co-organizer



User Scenarios

Multipart File Upload

Select files


Choose Files

No file chosen

Upload

User Scenarios

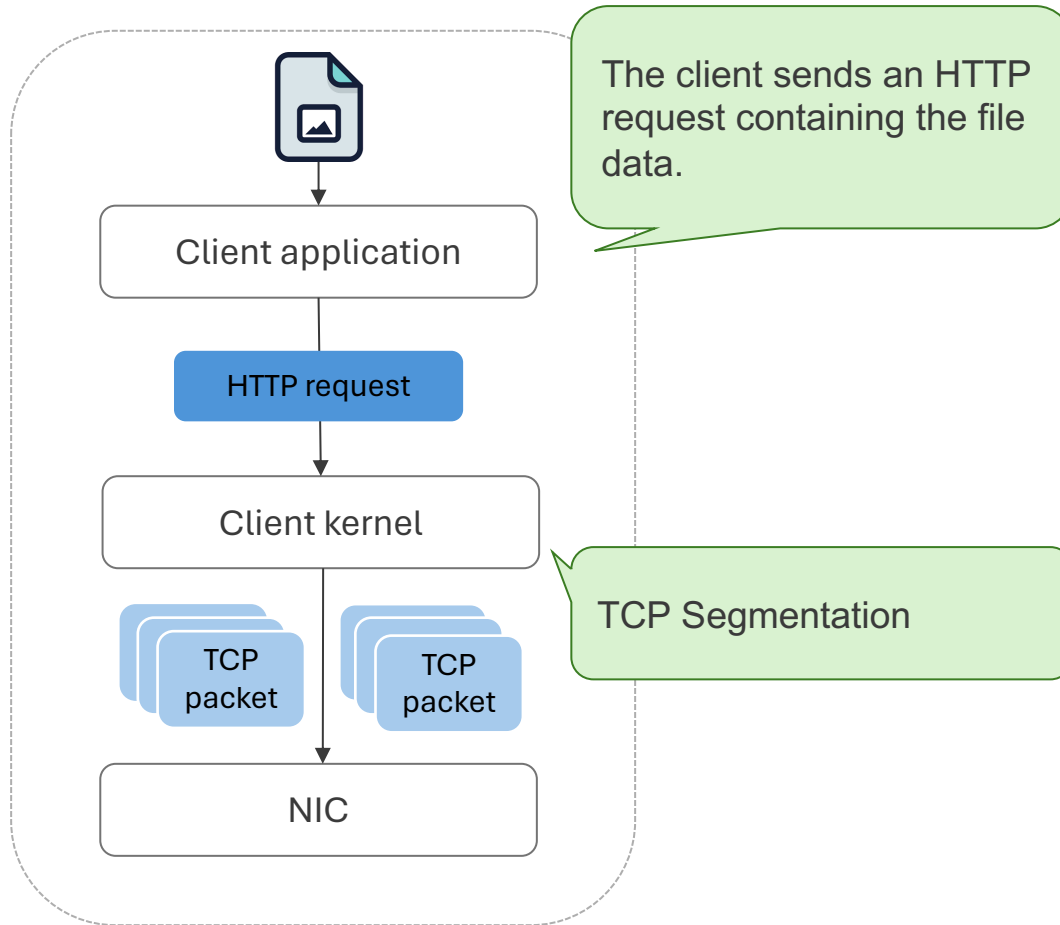
- A file uploading system that helps user upload their files
- Using HTTP/1.1 protocol
- Sync API



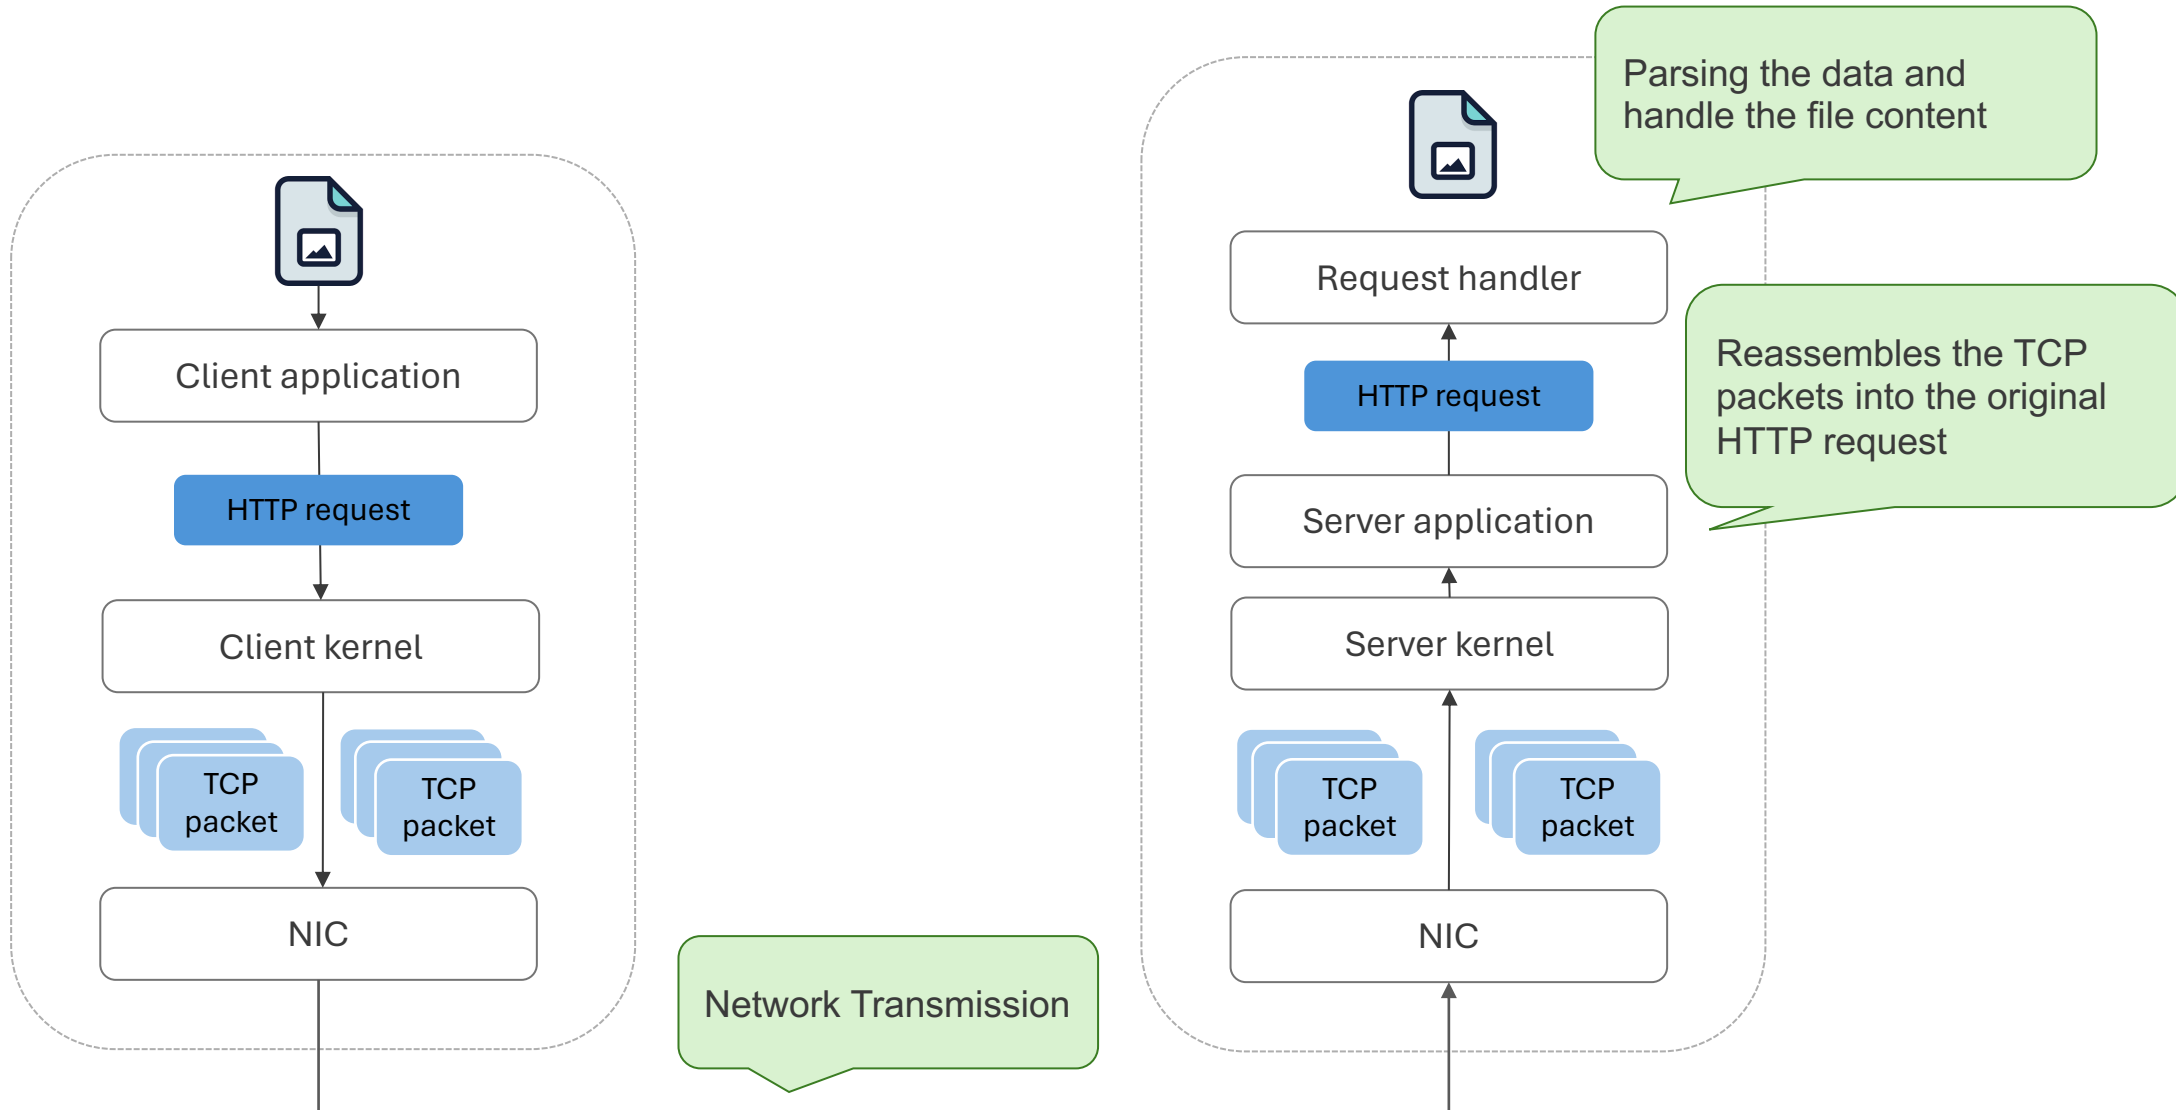
How to Handle Large File Uploads Efficiently

under the constraints

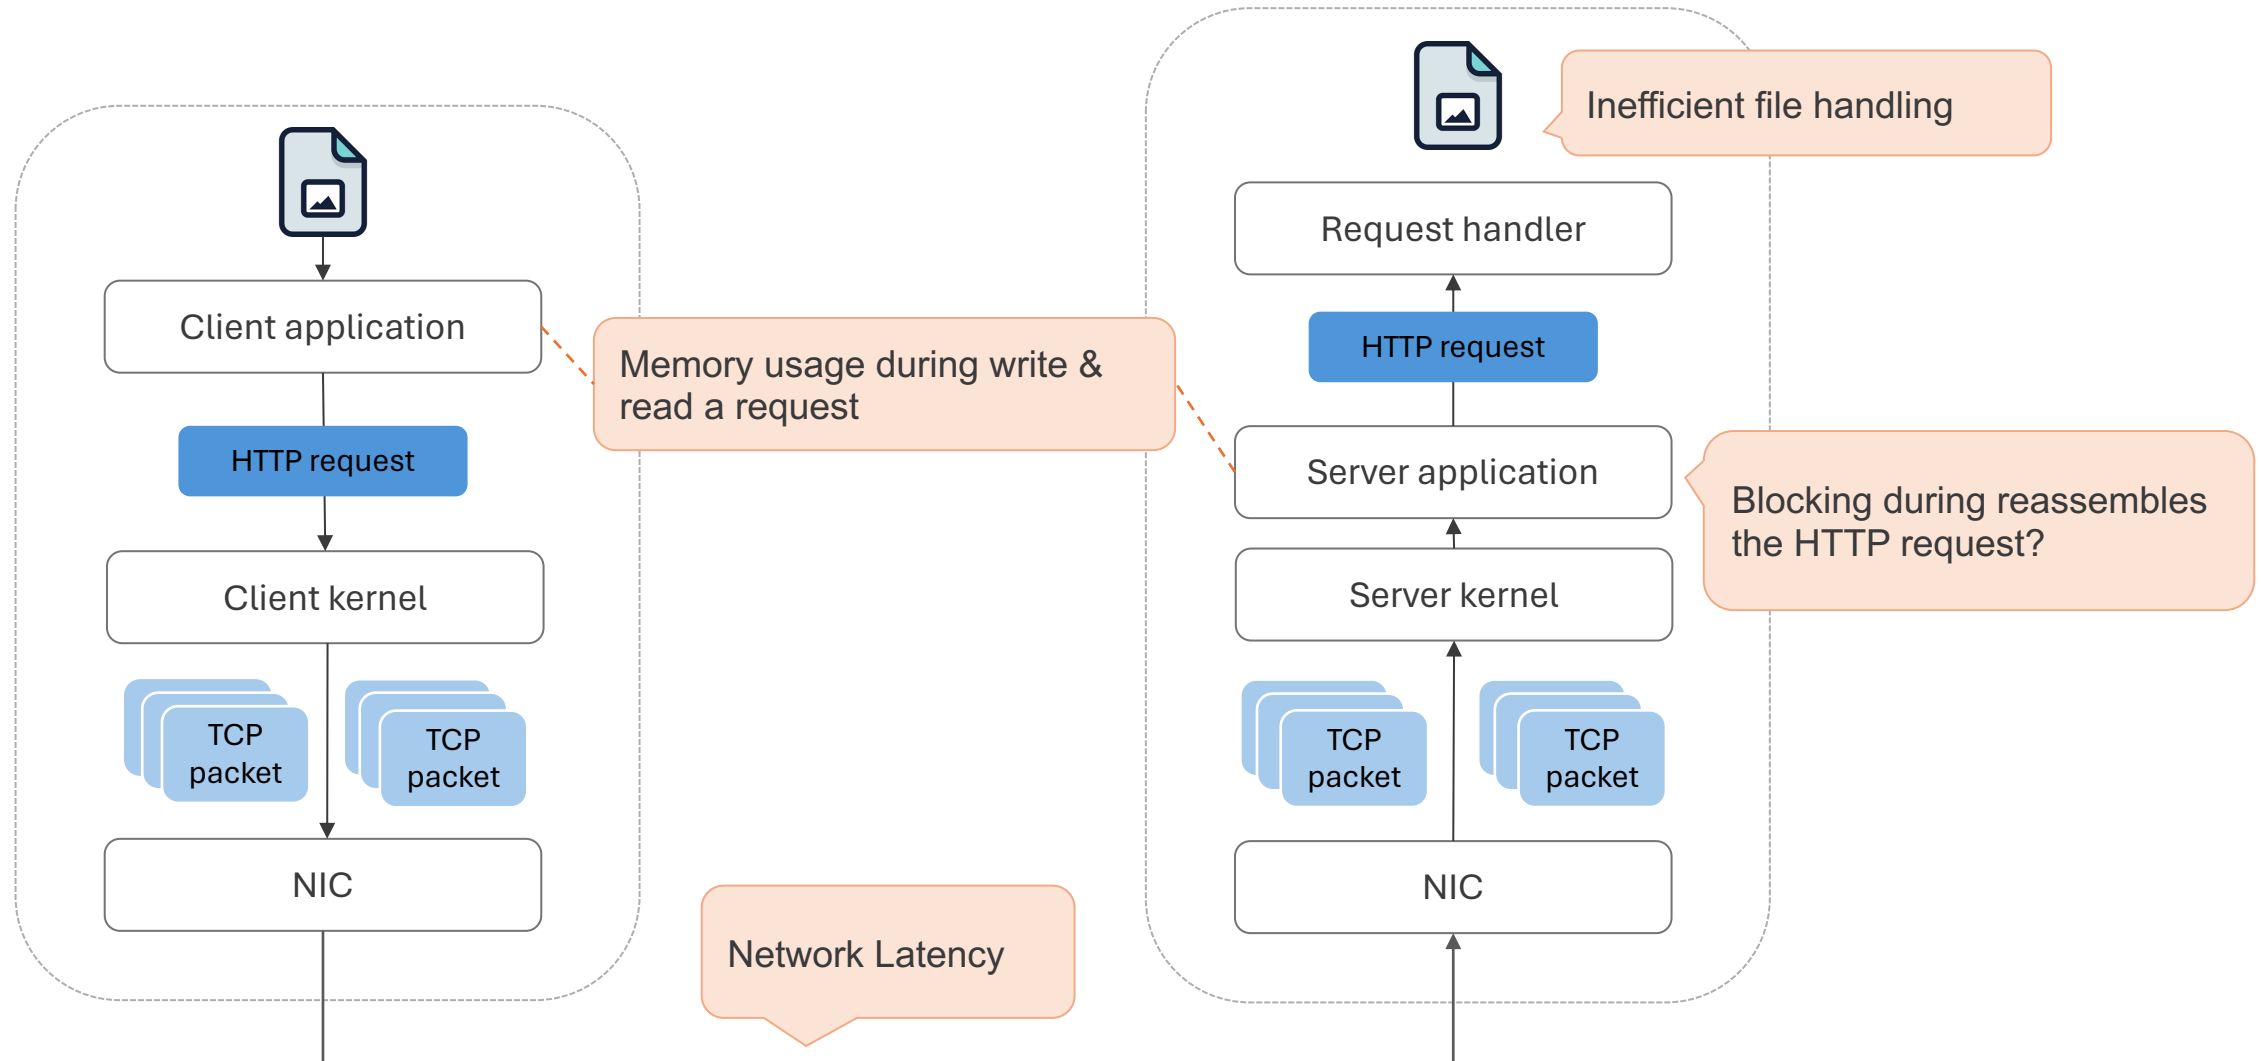
A Journey of a File Uploading Request



A Journey of a File Uploading Request



Risks



How Golang net/http Package Handle the Request

How Golang net/http package handle the request

HTTP Request Example

```
POST /upload HTTP/1.1
Host: golangtw.com
Content-Type: multipart/form-data; boundary=-----974767299852498929531610575
Content-Length: 554
```

Request headers

```
-----974767299852498929531610575
Content-Disposition: form-data; name="name"
John Doe
-----974767299852498929531610575
Content-Disposition: form-data; name="email"
johndoe@example.com
-----974767299852498929531610575
Content-Disposition: form-data; name="file";filename="profile.png"
Content-Type: image/png [...binary data of profile.png image file...]
-----974767299852498929531610575--
```

Part

Part

Part

Request body

How Golang net/http Package Handle the Request

```
POST /upload HTTP/1.1
Host: golangtw.com
Content-Type: multipart/form-data; boundary=-----974767299852498929531610575
Content-Length: 554
```

```
-----974767299852498929531610575
Content-Disposition: form-data; name="name"
John Doe
-----974767299852498929531610575
Content-Disposition: form-data; name="email"
johndoe@example.com
-----974767299852498929531610575
Content-Disposition: form-data; name="file";filename="profile.png"
Content-Type: image/png [...binary data of profile.png image file...]
-----974767299852498929531610575--
```

Parse the request and pass it to the handler

```
func handleUpload(w http.ResponseWriter, r *http.Request) {
    if err := r.ParseMultipartForm(maxMemory); err != nil {
        http.Error(w, "failed to parse form data", http.StatusBadRequest)
        return
    }

    file, _, err := r.FormFile(fileKey)
    if err != nil {
        http.Error(w, "failed to retrieve file", http.StatusBadRequest)
        return
    }
    defer file.Close()

    fileBytes, err := io.ReadAll(file)
    if err != nil {
        http.Error(w, "failed to read the file", http.StatusInternalServerError)
        return
    }
    // more code...
    slog.Info("successfully")
}
```

Parse the parts

Access the uploaded file

Question 1

Does net/http read the all content before passing the request to user's handlers?

```
POST /upload HTTP/1.1
Host: golangtw.com
Content-Type: multipart/form-data; boundary=-----974767299852498929531610575
Content-Length: 554
-----974767299852498929531610575
Content-Disposition: form-data; name="name"
John Doe
-----974767299852498929531610575--
```

The Go net/http package is responsible for **parsing the header and identifying the presence of a body.**

It doesn't automatically read the entire request body.

Question 2

What happens if the size of the content being uploaded is large?

```
func handleUpload(w http.ResponseWriter, r *http.Request) {  
    if err := r.ParseMultipartForm(maxMemory); err != nil {  
        http.Error(w, "failed to parse form data", http.StatusBadRequest)  
        return  
    }  
}
```

A file part exceeds **maxMemory**, it's written to a temporary file on disk.


Smaller file parts might be kept entirely in memory.

Default: **32 MB**

Question 2

What happens if the size of the content being uploaded is large?

```
func handleUpload(w http.ResponseWriter, r *http.Request) {  
  
    if err := r.ParseMultipartForm(maxMemory); err != nil {  
        http.Error(w, "failed to parse form data", http.StatusBadRequest)  
        return  
    }  
  
    file, _, err := r.FormFile(fileKey)   
  
    if err != nil {  
        http.Error(w, "failed to retrieve the file", http.StatusBadRequest)  
        return  
    }  
}
```



Risks

```
func handleUpload(w http.ResponseWriter, r *http.Request) {  
    if err := r.ParseMultipartForm(maxMemory); err != nil {  
        http.Error(w, "failed to parse form data", http.StatusBadRequest)  
        return  
    }  
  
    file, _, err := r.FormFile(fileKey)  
    if err != nil {  
        http.Error(w, "failed to retrieve the file", http.StatusBadRequest)  
        return  
    }  
    defer file.Close()  
  
    fileBytes, err := io.ReadAll(file)  
  
    if err != nil {  
        http.Error(w, "failed to read the file", http.StatusInternalServerError)  
        return  
    }  
    // more code...  
    slog.Info("successfully")  
}
```

handling many requests concurrently

many parts or complex structures within the form data

excessive memory usage

slow disk writes

Streaming Data: Chunked Request

Streaming Data: Chunked Request HTTP/1.1

POST /upload HTTP/1.1

Host: golangtw.com

Transfer-Encoding: chunked

Content-Disposition: attachment; filename="golang_is_good.txt"

Content-Length header is not required

Request 1

2000\r\n

[8192 bytes of file content]\r\n

2000\r\n

[8192 bytes of file content]\r\n

Request 2

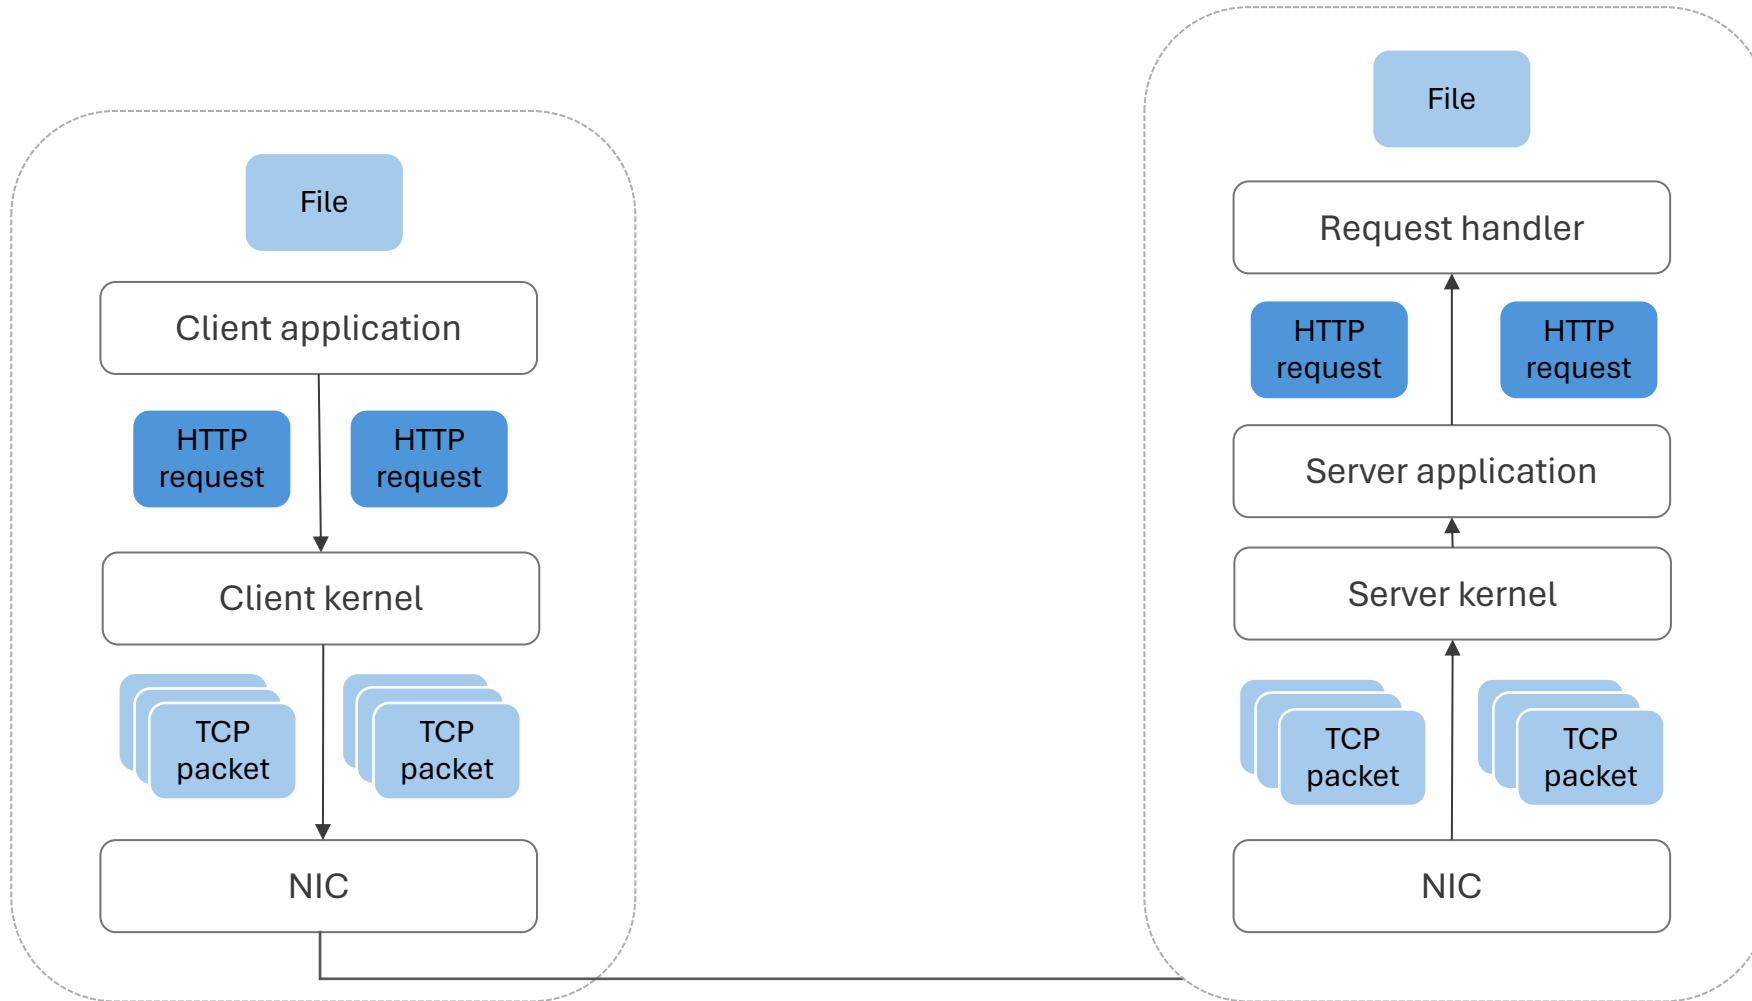
1500\r\n

[6144 bytes of file content]\r\n

0\r\n // Final chunk with size 0

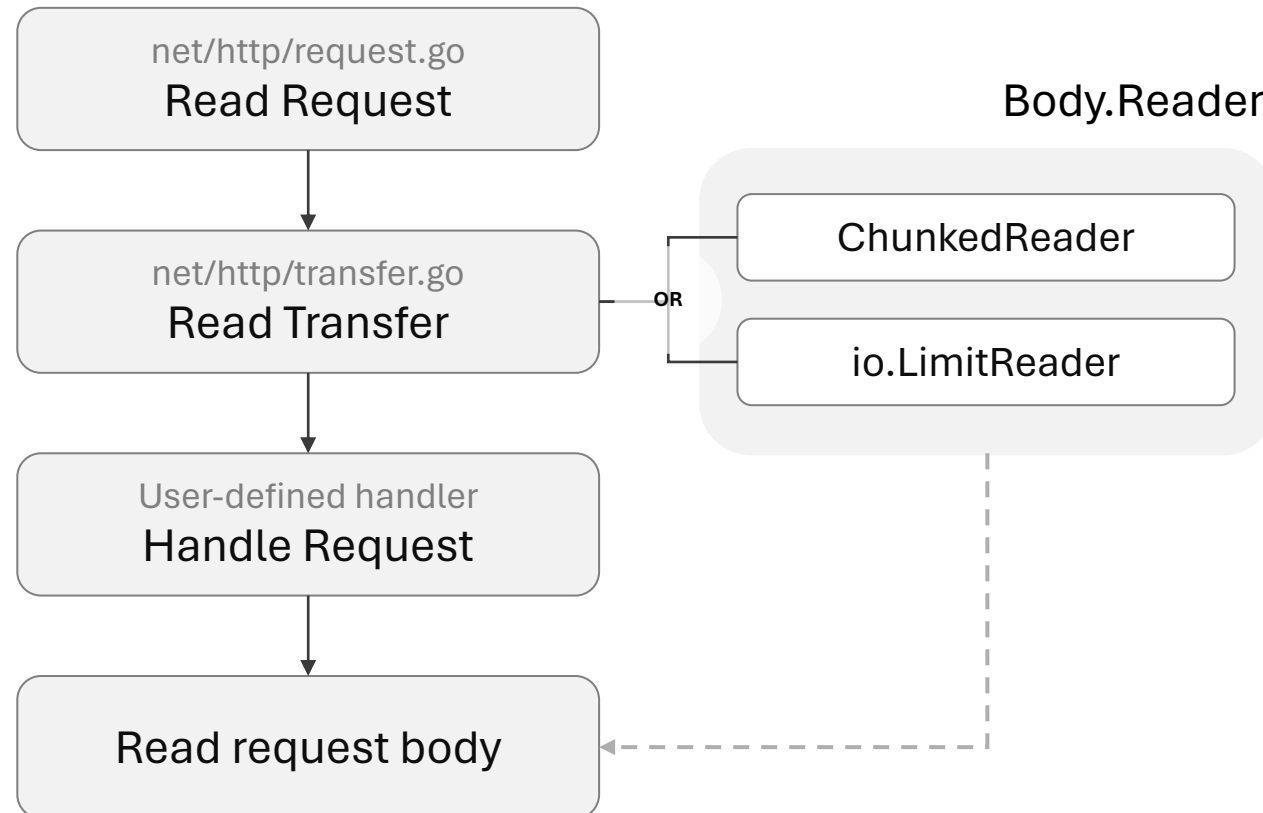
Request 3

Streaming Data: Chunked Request HTTP/1.1



How Golang net/http package handle chunked requests

The server receives a request with Transfer-Encoding:chunked, it **automatically dechunks** the data for user. User can access the dechunked data through the `http.Request.Body` field.



How can we optimize the code for better performance

How can we optimize the code
for better performance

Server Side - Original



excessive memory usage

```
const (  
    maxMemory      = 4096  
    fileKey        = "file"  
    uploadDestFolder = "uploads"  
)  
  
func handleUpload(w http.ResponseWriter, r *http.Request) {  
    if err := r.ParseMultipartForm(maxMemory); err != nil {  
        http.Error(w, "failed to parse form data", http.StatusBadRequest)  
        return  
    }  
  
    file, _, err := r.FormFile(fileKey)  
    if err != nil {  
        http.Error(w, "failed to retrieve the file", http.StatusBadRequest)  
        return  
    }  
    defer file.Close()  
  
    fileBytes, err := io.ReadAll(file)  
    if err != nil {  
        http.Error(w, "failed to read the file", http.StatusInternalServerError)  
        return  
    }  
  
    filePath := filepath.Join(uploadDestFolder, uuid.New().String())  
    err = os.WriteFile(filePath, fileBytes, 0644)  
    if err != nil {  
        http.Error(w, "failed to save the file", http.StatusInternalServerError)  
        return  
    }  
  
    slog.Info("File uploaded successfully to", slog.String("path", filePath))  
}
```

How can we optimize the code for better performance

Server Side - Limit Read

Benchmark	Average Time per Operation (ns)	Average Memory Allocation (B)
ReadAll	350,250,692	615,239,886
LimitRead	124,428,701	7,316



Limit Read the source content into the destination

```
func handleUpload(w http.ResponseWriter, r *http.Request) {
    if err := r.ParseMultipartForm(maxMemory); err != nil {
        http.Error(w, "failed to parse form data", http.StatusBadRequest)
        return
    }

    file, _, err := r.FormFile(fileKey)
    if err != nil {
        http.Error(w, "failed to retrieve the file", http.StatusBadRequest)
        return
    }
    defer file.Close()

    filePath := filepath.Join(uploadDestFolder, uuid.New().String())
    dstFile, err := os.Create(filePath)
    if err != nil {
        http.Error(w, "failed to create file", http.StatusInternalServerError)
        return
    }
    defer dstFile.Close()

    _, err = io.Copy(dstFile, file)
    if err != nil {
        http.Error(w, "failed to save the file", http.StatusInternalServerError)
        return
    }

    slog.Info("File uploaded successfully to", slog.String("path", filePath))
}
```

Default Buffer: 32KB

How can we optimize the code
for better performance

Server Side - Custom

Benchmark	Average Time per Operation (ns)	Average Memory Allocation (B)
ReadAll	350,250,692	615,239,886
LimitRead	124,428,701	7,316
Custom	89,152,902	5,301



Avoid saving the content into temp. file

```
func ParseMultipartFormAndSaveFile(r *http.Request, filePath string) error {  
    contentType := r.Header.Get("Content-Type")  
    if contentType == "" {  
        return fmt.Errorf("missing Content-Type header")  
    }  
    mediaType, params, err := mime.ParseMediaType(contentType)  
    if err != nil {  
        return fmt.Errorf("invalid Content-Type: %v", err)  
    }  
    if mediaType != "multipart/form-data" {  
        return fmt.Errorf("Content-Type is not multipart/form-data")  
    }  
  
    boundary := params["boundary"]  
    reader := multipart.NewReader(r.Body, boundary)  
  
    for {  
        part, err := reader.NextPart()  
        if err == io.EOF {  
            break  
        }  
  
        // more code  
  
        if part.FileName() != "" {  
            dst, err := os.Create(filePath)  
            if err != nil {  
                return fmt.Errorf("failed to create file: %v", err)  
            }  
            defer dst.Close()  
  
            if _, err := io.Copy(dst, part); err != nil {  
                return fmt.Errorf("failed to copy file content: %v", err)  
            }  
            return nil  
        }  
    }  
}
```

How can we optimize the code
for better performance

Client Side - Original

```
file, err := os.Open(filePath)
if err != nil {
    fmt.Printf("failed to open file: %v\n", err)
    return
}
defer file.Close()

var requestBody bytes.Buffer
writer := multipart.NewWriter(&requestBody)

formFile, err := writer.CreateFormFile("file", filepath.Base(file.Name()))
if err != nil {
    fmt.Printf("failed to create form file field: %v\n", err)
    return
}

_, err = io.Copy(formFile, file)
if err != nil {
    fmt.Printf("failed to copy file content: %v\n", err)
    return
}
writer.Close()

request, err := http.NewRequest("POST", url, &requestBody)
if err != nil {
    fmt.Printf("failed to create request: %v\n", err)
    return
}
request.Header.Set("Content-Type", writer.FormDataContentType())
client := &http.Client{}
response, err := client.Do(request)
```


How can we optimize the code for better performance

Client Side - Pre-Allocate



Pre-allocate the buffer to prevent slice.extend

```
file, _ := os.Open(filePath)
defer file.Close()

buf := make([]byte, 0, 157286400)
body := bytes.NewBuffer(buf)
writer := multipart.NewWriter(body)

h := make(textproto.MIMEHeader)
h.Set("Content-Disposition",
    fmt.Sprintf(`form-data; name="%s"; filename="%s`",
        escapeQuotes("file"), escapeQuotes(filepath.Base(file.Name()))))
h.Set("Content-Type", "application/vnd.ms-excel")

part, _ := writer.CreatePart(h)
io.Copy(part, file)
writer.Close()

header := http.Header{}
header.Add("Content-Type", writer.FormDataContentType())
```

How can we optimize the code for better performance

Client Side - Pipe



Using io.Pipe to prevent load the entire file into memory.

```
pipeReader, pipeWriter := io.Pipe()
writer := multipart.NewWriter(pipeWriter)

go func() {
    defer pipeWriter.Close()
    defer writer.Close()

    formFile, err := writer.CreateFormFile("file", filepath.Base(file.Name()))
    if err != nil {
        fmt.Printf("failed to create form file field: %v\n", err)
        return
    }

    _, err = io.Copy(formFile, file)
    if err != nil {
        fmt.Printf("failed to copy file content: %v\n", err)
        return
    }
}()

request, err := http.NewRequest("POST", url, pipeReader)
```

```
if err != nil {
    fmt.Printf("failed to create request: %v\n", err)
    return
}
```

Security Issue

Memory Exhaustion in Request.ParseMultipartForm

CVE-2023-45290 Detail

AWAITING ANALYSIS

This vulnerability is currently awaiting analysis.

Description

When parsing a multipart form (either explicitly with `Request.ParseMultipartForm` or implicitly with `Request.FormValue`, `Request.PostFormValue`, or `Request.FormFile`), limits on the total size of the parsed form were not applied to the memory consumed while reading a single form line. This permits a maliciously crafted input containing very long lines to cause allocation of arbitrarily large amounts of memory, potentially leading to memory exhaustion. With fix, the `ParseMultipartForm` function now correctly limits the maximum size of form lines.

QUICK INFO

CVE Dictionary Entry:

[CVE-2023-45290](#)

NVD Published Date:

03/05/2024

NVD Last Modified:

05/01/2024

Source:

Go Project

Fixed: 2024/05/06

Wrapping Up

- How the Golang net/http package handles HTTP requests
 - Passes requests to user-defined handlers without reading the entire request body
 - May store form data in temporary files if the size exceeds a limit
 - Automatically handles chunked requests
- Performance improvement in server-side
 - Use limit reading method to read the request body
 - Customizing the file-handling process to avoid saving the content to a temporary file
- Performance improvement in client-side
 - Using io.Pipe to prevent load the entire file into memory